

Deploying Liferay DXP Using Docker

Table of Contents

Introduction	1
What Is Docker?	1
Shipping Liferay DXP Using Docker	2
Building Docker Images for Liferay Components. . .	2
Liferay DXP	3
JDK.	3
Recommended: OpenJDK.	3
Alternative: Oracle JDK.	3
Liferay Bundle.	4
Exposing Liferay Ports	5
Starting Application Server	5
LCS and Liferay DXP Activation Keys	5
Other Considerations	6
Elasticsearch	7
Options for Liferay Document Storage	8
Docker Volumes and Volume Plugins.	8
Alternate Store Implementations	8
Options for Mounting Configurations Into Docker. .	9
Options for Configuring Networking.	9
Docker Networking	10
Legacy Container Links.	11
Summary	11
Moving Forward.	11
Liferay Global Services	11
Further Reading	11

Introduction

There has been growing interest in using Liferay Digital Experience Platform (DXP) with Docker, a popular software container platform that makes it easier to use containers to create, run, test and deploy applications faster. This guide aims to provide a starting point for companies that would like to use Docker as a runtime environment for Liferay DXP. It will also provide guidance for how a basic Liferay DXP deployment can be modeled using Docker's capabilities.

We will cover several key topics:

- How to create a basic Docker image for Liferay DXP and Elasticsearch
- Options for document storage, including Docker volume plugins and Document implementations not based on a file system
- How to mount configurations into Docker
- How to configure networking

What Is Docker?

[Docker](#) is described by its creators as follows:

Docker allows you to package an application with all of its dependencies into a standardized unit for software development.

These units are called Docker images and they can be run as containers in supported operating systems. The image can define one or more network ports that will be exposed from the container (accessible from the network) and other containers or standalone applications. It can then use these ports to interact with the application running inside the container.

Docker is a virtualization concept similar to virtual machines, but instead of trying to mimic physical hardware as closely as possible (which is the aim of traditional virtualization providers like VirtualBox or VMWare), Docker's main goals are:

- Being descriptive in how the Docker image was built. Every image is described as a set of individual configuration steps executed on top of the base image. These steps can be written in the form of a [Dockerfile](#).
- Keeping images smaller than a full-featured virtual machine. This is possible thanks to the fact that Docker containers are sharing the kernel with the host OS. Necessary tools, libraries and/or applications can be installed as needed.

- Quicker startup/shutdown of containers compared to virtual machines. Also, the guest OS used with Docker is often a stripped-down version of the OS, leaving only what's crucial, which helps the speed even more.

Docker containers run natively in all modern Linux systems and supplemental software is available to run containers in OS X- or Windows-based systems.

Docker was released in early 2013 and has gained huge popularity, especially in the microservices-oriented world and among DevOps. It has both supporters and [critics](#), and the use of Docker, especially in production, should be carefully considered.

Shipping Liferay DXP Using Docker

Liferay DXP can operate in a variety of environments and can be configured to run inside a Docker container as well. Liferay DXP is a standard Java-based web application that can run on many application servers and supported operating systems. Java application servers are Java Virtual Machine (JVM) based; success is defined by hosting a JVM with application server binaries inside and Liferay DXP deployed.

Liferay DXP can therefore run in any Docker container that has the JVM installed. Running DXP inside a Docker container is functionally identical to running a standard Linux distribution.

Please note that in order to be able to receive support from Liferay, you will need to select a supported Application server, database, JDK and operating system. Please check Liferay's [Compatibility Matrix](#) for more details.

Building Docker Images for Liferay Components

Liferay DXP has two major components which need computational power and therefore are good candidates for becoming separate Docker containers:

- Liferay DXP itself, including the front end and back end, processing all requests coming to Liferay DXP
- The search engine, as a separate system where indexing and searching of Liferay assets is performed

Liferay also needs to store its data persistently, which is handled by:

- Access to an existing relational database, either running in another Docker container or as a traditional database server listening on a known hostname and port. This topic won't be covered in this guide. Please see other Liferay documentation on setting up a database for Liferay DXP.
- Access to location for storing Documents & Media files. This will be covered in more detail later in this guide.

Liferay DXP

To build a Docker image with Liferay DXP, the following high-level steps will need to be accomplished:

- Install supported JDK into the image
- Install Liferay bundle into the image
- Expose application server ports (HTTP, AJP, etc.)
- Start application server with Liferay

JDK

Liferay DXP should be used with one of [the supported VMs](#), which is currently either Oracle JDK 8, IBM JDK or OpenJDK.

RECOMMENDED: OPENJDK

OpenJDK will be the best choice with Docker for two reasons:

1. You can easily have it preinstalled by basing your image on Docker image *openjdk:8-jre*.
2. If you want to use external Elasticsearch nodes — the only supported option in Liferay DXP — you have to keep the VMs identical (due to TCP channel being used between Liferay DXP and Elasticsearch) and official Elasticsearch images use OpenJDK.

ALTERNATIVE: ORACLE JDK

If you choose to use Oracle JDK, it can be installed into the image in a variety of ways:

1. Download the JDK from Oracle's website manually and include it in the image.
 - a. Go to [oracle.com](https://www.oracle.com) and find the JDK distribution appropriate to the Docker image kernel.

- b. ADD or COPY instructions in a Dockerfile can be used to include/install the JDK.
2. Download the JDK as part of the build process.
3. Use scripts to download and install the JDK when the Docker image is built.
 - a. Use RUN instructions and, for example, curl / wget.
4. Use another method.
 - a. Use custom PPA in Ubuntu if it is suitable for the kernel used in the Docker image.

Please note that using any automated approach of downloading the JDK (options two through four above) implies agreement with the *Oracle Binary Code License Agreement for the Java SE Platform Products and JavaFX*. The JDK cannot be downloaded without the user accepting the license agreement, so any automated method has to contain the agreement as part of the download, expressed by presenting an HTTP cookie in the issued HTTP request. By using, for example, the script, you agree to the terms and conditions just as if you had downloaded the JDK manually (option one above).

Please also note that redistributing Oracle JDK is limited by Oracle's licensing terms. In short, never *publicly* distribute a Docker image with Oracle JDK installed unless the image:

- Also contains the application using the JDK (e.g., Liferay DXP) and
- The JDK is complete — none of the original files were deleted or altered after installation, for example to minimize the size of the image.

If the JDK files are altered or the application isn't bundled with the JDK, the docker image can only be shared internally. You are free to distribute it within your organization, but never publicly (e.g., on Docker Hub) or to clients or partner companies.

We recommend carefully reviewing Oracle's license agreement and [how it works with Docker](#) before using it with Liferay DXP.

LIFERAY BUNDLE

Liferay bundle files need to be installed into the Docker image. Extract the files from the Liferay DXP archive into the desired location in the Docker image, using ADD or COPY instructions in a Dockerfile, or eventually using RUN instructions to execute some bundle installation script.

EXPOSING LIFERAY PORTS

Use the EXPOSE instruction to instruct Docker which ports should be exported from the image, making them accessible from outside of the container. The exact port numbers depend on the setup of the application server.

In Tomcat, for example, port 8080 (HTTP connector) or 8009 (AJP connector) is exposed when using the example Tomcat setup.

If you want to expose the OSGi Gogo shell, expose port 11311 too, but take into consideration the implications of doing this in production. For this to work, you need to open Gogo console to any network interface: `module.framework.properties.osgi.console=0.0.0.0:11311` in order to be able to access it remotely.

STARTING APPLICATION SERVER

As in the previous step, this is highly dependent on the setup. In general, use the ENTRYPOINT instruction to execute a startup script shipped with the application server.

As an example, using Tomcat 8 and Liferay DXP home `/opt/liferay/liferay-portal-tomcat`, the command would be: `/opt/liferay/liferay-portal-tomcat/tomcat-8.0.32/bin/catalina.sh run`. This assumes Tomcat is located in a subdirectory named `tomcat-8.0.32`.

LCS AND LIFERAY DXP ACTIVATION KEYS

If you are using Liferay Connected Services (LCS) to provision activation keys from your Liferay DXP subscription, you need to be aware of the process by which the activation keys are provided and released and how the Docker containers are terminating Liferay DXP running inside.

During the shutdown of Liferay DXP (i.e., the shutdown of its app server, like Tomcat), LCS actively releases any activation keys that are currently in use. For that to happen, the process of the JVM cannot be simply killed in the OS. It has to be allowed to shut down. Unfortunately, this does not often happen in Docker, where typically `bin/catalina.sh run` is started in the foreground, using ENTRYPOINT or CMD instructions. In such cases, when the Docker container is killed, SIGTERM signal is sent to the process, ending its run immediately.

If the activation key is not released actively, it will be freed automatically if a given Liferay DXP node does not send a heartbeat signal to LCS for more than six minutes.

With your LCS subscription, you cannot consume any extra activation keys when all are currently used. That means you cannot get an activation key for any newly started Liferay DXP nodes until you release the keys used in your current

set of nodes. Your new set of Liferay DXP nodes will start up without receiving activation keys from LCS for roughly six minutes after they are started. You will only see the LCS subscription page when browsing Liferay DXP.

If you want to avoid this, you will need to implement a custom startup script for Liferay DXP in Docker, which traps all the termination signals and waits until the Java process of the JVM fully completes its shutdown. You can refer to the Knowledge Base on Liferay's Customer Portal for full details and a sample startup script handling the shutdown of Liferay DXP in Docker.

Note: Liferay is planning to reduce the six minute window to two minutes in a future version of LCS.

OTHER CONSIDERATIONS

Although not mandatory, it is good to follow common Liferay/Docker conventions:

- Consider creating a dedicated user inside Docker image for running the application server (i.e., create a user called `liferay` or `app`).
 - Although the security of Docker containers has improved over the years, if the application runs as `root` (inside the container), it can pose a possible risk to the host machine.
 - Take permissions into consideration. `COPY/ADD` instructions use `root` to copy files/folders into the image, so on runtime, it could be a problem if the user that manages the application cannot read/access/write those files. The usage of `chown` as a Dockerfile instruction can help, although it will increase the image size. The files affected by `chown` will be added to a new layer.
- Use the dedicated user to start Liferay DXP (application server), not `root`.
- Always start only one main process in the container — Liferay DXP. Other applications should be started in their own container or outside of Docker.

Elasticsearch

In Liferay DXP, the default search engine is an *embedded* Elasticsearch (ES) server. Although this approach is an easy out-of-the-box solution, Liferay will only support a search engine running *outside* of the JVM running Liferay DXP. These two applications need very different tuning parameters and many of the issues stem from sharing the JVM for processing requests and indexing/searching the data in the same JVM. External Solr server engine is supported the same

way as in previous versions of the Liferay platform. This document focuses on external Elasticsearch engine setup, but the principle of setting up Solr would be very similar.

Since Elasticsearch should run inside Docker, a separate image with its installation is needed. Although the Docker image can be created manually, for example by writing a Dockerfile, it will be more straightforward to use the official image for Elasticsearch which is [available on DockerHub](#). Moreover, due to the fact that Liferay DXP uses TCP to communicate with Elasticsearch, the VMs of these two applications should be as similar as possible, but at minimum:

- the same vendor (e.g., either Oracle JDK or OpenJDK);
- the same major version (JDK 8 in our case — JDK 7 is not supported in Liferay DXP).

Since OpenJDK is now supported for Liferay DXP, the best approach is to base both images on *openjdk:8-jre*.

Once a Docker container with an Elasticsearch node is started, Liferay DXP needs to be told to use the external Elasticsearch engine and which endpoint of Elasticsearch should be used for interacting with this engine. For the necessary Liferay DXP configuration, check Liferay documentation. Please note that only Elasticsearch 2.2.x instances are currently supported by Liferay DXP.

If deploying multiple Elasticsearch containers to form a cluster of Elasticsearch nodes, the situation becomes more complicated since Docker networking needs to be taken into consideration. Please see section [Options for Configuring Networking](#) for more details, which apply both to clustering Elasticsearch and Liferay DXP.

Options for Liferay Document Storage

By default, Liferay DXP stores *Documents & Media* files (their binary content) on the local file system. With default configuration (`FileSystemStore` or `AdvancedFileSystemStore` being used), this will be inside `data/document_library` directory in Liferay DXP's home directory, unless changed using OSGi configurations (i.e., `data/my_own_file_system_structure`). Since all file system data inside the Docker container is ephemeral and available only for the time of the container's life, the data needs to be stored outside of the container to be persistent. There are several options to achieve this task.

Docker Volumes and Volume Plugins

One option is to leverage *Docker volumes*. On the container's start, the volume will be mounted from the container to the host, possibly to a specific mount point on the host. This will make sure that any change made in the mounted folder from inside the container will be visible in the directory on the host and vice versa. A Docker volume can be defined by adding a `VOLUME` directive into the Dockerfile. This solution is, however, host-dependent and requires a setup outside of the Docker image to share the volumes between multiple Containers or even Docker hosts.

If running a cluster of Liferay DXP nodes, the *Documents & Media* directory needs to be shared between all the nodes. Docker offers support to enable Docker volumes to persist across multiple Docker hosts, using *Volume plugins*. Each of these plugins offers different implementation on how the defined volumes can be shared between multiple containers and hosts. Please check Docker documentation for more details.

This principle is similar to the use of Network-Attached Storage (NAS) devices, SAN drives or NFS-mounted volumes in a non-Docker deployment of Liferay DXP.

Alternate Store Implementations

Another option is to use an implementation of Liferay store which does not use the local file system. That way, the data is not stored locally in the Docker container and even when it ceases to exist, the data still exists in the outside system.

Some examples could be `S3Store`, storing data in a configured S3 bucket inside Amazon Web Services, or `DBStore`, storing binary files in Liferay's relational database. Please note that each of the stores has its usage specifics and `DBStore` is not recommended for use in production environments with moderate or heavy use of *Documents & Media*. Please see Liferay documentation for further details.

Options for Mounting Configurations Into Docker

In Liferay DXP, the configuration is read from three basic places:

1. The legacy properties in *portal.properties* and its overrides using various *portal-*.properties*.

2. Since [Liferay began to support reading configuration properties from an OS environment](#), it is possible to set up environment variables into the machine hosting Liferay, following the name convention described in the linked JIRA ticket. The Setup Wizard could be set in this way: `LIFERAY_SETUP_PERIOD_WIZARD_PERIOD_ENABLED=false`.
3. Files loaded by *OSGi Config Admin*, placed inside `[lifera_y_home]/osgi/configs`.

One option for configuring an existing Liferay bundle inside a Docker image is by identifying the Docker volumes that can be used to propagate files and folders, and mounting them from the host into the desired location in the Liferay bundle inside of the container. Check Docker volumes documentation for more details.

Another option is to set up a configuration system (such as *etcd*, *Consul* or *ZooKeeper*) to hold the important configuration keys and values, and then configure local daemon (such as *confd*) to read the configuration and write the necessary files into the Liferay DXP bundle inside the container.

Please note that for production deployments, Liferay recommends placing all relevant configurations into the Liferay bundle directly and including it in the Docker image as a whole, instead of mounting additional configurations into a container and possibly changing the files during Liferay's runtime. It might, however, be useful during the development phase, when testing various configuration changes (assuming Liferay DXP supports the reloading of a given type of configuration from the file).

Options for Configuring Networking

One of the basic networking needs of Liferay DXP is the exposure of the app server ports of the HTTP or AJP connectors, which is straightforward in Docker as shown in the [Exposing Liferay ports](#) section. However, there are some additional steps if you plan to run Liferay DXP in a clustered setup, since the operation of Cluster Link networking will also need to be configured.

With default configuration, Cluster Link relies on multicast for nodes discovery and UDP for communication between cluster nodes. This setup might not be reliable in virtualized environments, like Docker, so it's safer to opt for using alternate discovery protocols (JDBC_PING or FILE_PING) and TCP unicast for communication. Since FILE_PING would require a shared file system between all Liferay node containers, it's usually easier to set up JDBC_PING, requiring similar configuration as configuring relational database access for Liferay DXP itself.

Once the discovery of Liferay nodes is handled, the jGroups' ports must be opened for incoming network traffic and the bind address of jGroups must be resolvable to an IP that is reachable by all other Docker containers (other Liferay nodes running in Docker containers) within the defined Docker network. Cluster Link uses the jGroups library for communication between the nodes. It typically needs two ports to bind to — one for control channel, one for transport channel.

Docker Networking

Docker added [container networking capabilities](#) into the base product, and it is now supported natively by the Docker engine itself. Several network drivers are available and are very well described in official Docker documentation.

Networking on one Docker host (all containers running on the same machine) is typically not problematic and can be handled using networks with the *bridge* driver. Use either the default *bridge* network (named the same as the underlying driver it uses) or create a custom one with the *bridge* driver. All Docker containers started into this network can immediately communicate with each other. A private network is created and each new container is given a new address within this network, just as a DHCP client would do in a standard networking infrastructure.

Multi-host networking is more involved. The *bridge* network driver does not support this type of networking. Instead, an *overlay* driver needs to be used. This driver requires an external service being available (a key-value store) and uses it to organize the routing of the traffic between multiple Docker hosts. After being set up, the container will have the same logical view of the network, even though it is physically spanning over multiple Docker hosts. Please check Docker documentation for detailed information.

Docker also supports the installation and use of *Networking plugins* which offer additional features and possibilities for how to make networking between Docker containers work. See Docker documentation for more details.

Legacy Container Links

Before explicit networking support in Docker was implemented, users could use Container links. Docker containers could “import” ports exposed by another running Docker container by aliasing them into the container, including the aliases for the hostnames.

Please note that using container links is discouraged in Docker's documentation and only works on the default bridge network, for backward compatibility.

Summary

This document outlines how Liferay DXP can be deployed and run as a Docker container, together with an external Elasticsearch in a separate Docker container. The necessary steps were documented to build a Docker image with a Liferay DXP bundle and Oracle JDK installed inside and how to run a Docker container based on this image.

Some advanced topics were also discussed, like options for Document & Media storage within Docker infrastructure and possible ways to plug configuration into a Docker container with Liferay DXP. Important considerations to take into account when defining and implementing networking between Liferay Docker nodes were listed.

Moving Forward

Liferay Global Services

Learn how Liferay's Global Services team can support your Liferay DXP project. Our consultants work alongside your team to create a personalized solution for your business. Contact us at liferay.com/consulting

Further Reading

- “Docker - Build, Ship, and Run Any App, Anywhere”: <https://www.docker.com>
- “Docker on Wikipedia”: [https://en.wikipedia.org/wiki/Docker_\(software\)](https://en.wikipedia.org/wiki/Docker_(software))
- “A brief history of Docker Containers’ overnight success”: <http://searchservervirtualization.techtarget.com/feature/A-brief-history-of-Docker-Containers-overnight-success>
- “Docker’s just a bit dodgy, but ready for rollout says Gartner”: http://www.theregister.co.uk/2015/01/12/docker_security_immature_but_not_scary_says_gartner/
- “Documentation Redistribution Policy”: <http://www.oracle.com/technetwork/java/redis-137594.html>
- “Running Java on Docker? You’re Breaking the Law”: <http://blog.takipi.com/running-java-on-docker-youre-breaking-the-law/>
- “Java SE Development Kit 8 Downloads”: <http://www.oracle.com/technetwork/java/javase/downloads/jdk8-downloads-2133151.html>

- “Downloading Java JDK on Linux via wget is shown license page instead”: <https://stackoverflow.com/questions/10268583/downloading-java-jdk-on-linux-via-wget-is-shown-license-page-instead>
- “Install Oracle Java 8 in Ubuntu or Linux Mint via PPA repository [jdk8]”: <http://www.webupd8.org/2012/09/install-oracle-java-8-in-ubuntu-via-ppa.html>
- “Oracle Binary Code License Agreement for the Java SE Platform Products and JavaFX”: <http://www.oracle.com/technetwork/java/javase/terms/license/index.html>
- “Oracle Java SE and Oracle Java Embedded Products”: <http://www.oracle.com/technetwork/java/javase/terms/products/index.html>
- “Five security concerns when using Docker”: <https://www.oreilly.com/ideas/five-security-concerns-when-using-docker>
- “Manage data in containers”: <https://docs.docker.com/engine/userguide/containers/dockervolumes/>
- “Understand Engine plugins: Volume plugins”: <https://docs.docker.com/engine/extend/plugins/#volume-plugins>
- “Understand Docker container networks”: <https://docs.docker.com/engine/userguide/networking/dockernetworks/>
- “Legacy container links”: https://docs.docker.com/engine/userguide/networking/default_network/dockerlinks/
- “Understand Engine plugins: Network plugins”: <https://docs.docker.com/engine/extend/plugins/#network-plugins>



Liferay makes software that helps companies create digital experiences on web, mobile and connected devices. Our platform is open source, which makes it more reliable, innovative and secure. We try to leave a positive mark on the world through business and technology. Hundreds of organizations in financial services, healthcare, government, insurance, retail, manufacturing and multiple other industries use Liferay. Visit us at liferay.com.

© 2018 Liferay, Inc. All rights reserved.